

---

# Bandit networks

---

**Charles AUGUSTE**  
École des Ponts  
charles.auguste@ponts.org

**Louis TREZZINI**  
École des Ponts  
louis.trezzini@ponts.org

## Abstract

A single-agent multi-armed bandit (MAB) problem is a problem in which a gambler, being in face of several slot machines, has to decide which machines to play, how many times and in which order to play each machine. In this project, we consider a multi-agent MAB scenario in which a group of agents connected through a social network are engaged in playing a stochastic MAB game and focus on minimizing their regret. Every time a player takes an action, the reward is observed by both the agent and its neighbors in the network. The goal of this project is to understand different collaborative policies (NAIC-type and FYL) described in [1], and to compare their performance over different network structures.

## 1 Introduction

This project is a review of paper [1]. As such, we will only succinctly present their model and their main results, and refer to additional results in [1] when necessary.

The interesting part of bandit networks is that we consider multiple agents playing the same multi-armed bandit (MAB) problem. Instead of playing independently, they can be connected through a graph and share their knowledge of the game (actions taken and resulting rewards) with their neighbors. Moreover, they can decide to play very diverse policies, which can be much more than simple extensions of UCB or Thomson-Sampling to the network-setting. These collaborative policies can really profit of the structure of the graph considered and even adapt to the locally available information.

## 2 The model

We use the definitions and notations from [1]. Let's restate them quickly.

### 2.1 Single-agent stochastic multi-armed bandit (MAB) problem

Let  $\mathcal{K} = \{1, 2, \dots, K\}$  be the set of arms available to the agent. Each arm is associated with a distribution  $\mathcal{P}_k$ , and let  $\mu_k$  be the corresponding mean, unknown to the agent. Let  $n$  be the time horizon or the total number of rounds. In each round  $t$ , the agent chooses an arm, for which he receives a reward, an i.i.d. sample drawn from the chosen arm's distribution. The agent can use the knowledge of the chosen arms and the corresponding rewards upto round  $(t - 1)$  to select an arm in round  $t$ . The goal of the agent is to maximize the cumulative expected reward up to round  $n$ .

### 2.2 Multi-agent stochastic multi-armed bandit (MAB) problem

We consider a set of users  $V$  connected by an undirected fixed network  $G = (V, E)$ , with  $|V| = m$ . Assume that each user is learning the same stochastic MAB problem i.e., faces a choice in each time from among the same set of arms  $\mathcal{K}$ . In the  $t^{th}$  round, each user  $v$  chooses an arm, denoted

by  $a^v(t) \in \mathcal{K}$ , and receives a reward, denoted by  $X_{a^v(t)}^v(t)$ , an i.i.d. sample drawn from  $\mathcal{P}_{a^v(t)}$ . In the stochastic MAB problem set-up, for a given user  $v$ , the rewards from arm  $i$ , denoted by  $\{X_i^v(t) : t = 1, 2, \dots\}$ , are i.i.d. across rounds. Moreover, the rewards from distinct arms  $i$  and  $j$ ,  $X_i^v(t)$ ,  $X_j^v(s)$ , are independent. If multiple users choose the same action in a certain round, then each of them gets an independent reward sample drawn from the chosen arm's distribution. We use the subscripts  $i$ ,  $v$  and  $t$  for arms, nodes and time respectively. The information structure available to each user is as follows. A user  $v$  can observe the actions and the respective rewards of itself and its one hop neighbors in round  $t$ , before deciding the action for round  $(t + 1)$ .

The policy  $\Phi^v$  followed by a user prescribes actions at each time  $t$ ,  $\Phi^v(t) : H^v(t) \rightarrow \mathcal{K}$ , where  $H^v(t)$  is the information available with the user till round  $t$ . A policy of the network  $G$ , denoted by  $\Phi$ , comprises of the policies pertaining to all users in  $G$ . The performance of a policy is quantified by a real-valued random variable, called *regret*, defined as follows. The regret incurred by user  $v$  for using the policy  $\Phi^v$  up to round  $n$  is defined as,

$$R_{\Phi}^v(n) = \sum_{t=1}^n (\mu^* - \mu_{a^v(t)}) = n\mu^* - \sum_{t=1}^n \mu_{a^v(t)},$$

where  $a^v(t)$  is the action chosen by the policy  $\Phi^v$  at time  $t$ , and  $\mu^* = \max_{1 \leq i \leq K} \mu_i$ . We refer to the arm with the highest expected reward as the optimal arm. The regret of the entire network  $G$  under the policy  $\Phi$  is denoted by  $R_{\Phi}^G(n)$ , and is defined as the sum of the regrets of all users in  $G$ . The expected regret of the network is given by:

$$\mathbb{E}[R_{\Phi}^G(n)] = \sum_{v \in V} \sum_{i=1}^K \Delta_i \mathbb{E}[T_i^v(n)], \quad (1)$$

where  $\Delta_i = \mu^* - \mu_i$ , and  $T_i^v(n)$  is the number of times arm  $i$  has been chosen by  $\Phi^v$  up to round  $n$ . We omit  $\Phi$  from the regret notation, whenever the policy can be understood from the context. Our goal is to devise learning policies in order to minimise the expected regret of the network.

Let  $\mathcal{N}(v)$  denote the set consisting of the node  $v$  and its one-hop neighbours. Let  $m_i^v(t)$  be the number of times arm  $i$  has been chosen by node  $v$  and its one-hop neighbors till round  $t$ , and  $\hat{\mu}_{m_i^v(t)}$  be the average of the corresponding reward samples. These are given as:

$$m_i^v(t) = \sum_{u \in \mathcal{N}(v)} T_i^u(t)$$

$$\hat{\mu}_{m_i^v(t)} = \frac{1}{m_i^v(t)} \sum_{u \in \mathcal{N}(v)} \sum_{k=1}^t X_{a^u(k)}^u(k) \mathbb{I}\{a^u(k) = i\},$$

where  $\mathbb{I}$  denotes the indicator function. We use  $m_i^G(t)$  to denote the number of times arm  $i$  has been chosen by all nodes in the network till round  $t$ .

### 3 Policies

#### 3.1 UCB-Network

They propose the natural extension of the well-known single agent policy UCB to the network-setting, that they call "UCB-user". When each user in the network follows the UCB-user policy, the network policy is called UCB-Network and is outlined in Algorithm 1.

The following theorem presents an upper bound on the expected regret of a generic network, under the UCB-Network policy.

**Theorem 1** *Assume that the network  $G$  follows the UCB-Network policy to learn a stochastic MAB problem with  $K$  arms. Further, assume that the rewards lie in  $[0, 1]$ . Then,*

(i) *The expected total regret of  $G$  is upper bounded as:*

$$\mathbb{E}[R^G(n)] \leq b + \sum_{i: \mu_i < \mu^*} C_G \frac{8 \ln n}{\Delta_i},$$

---

**Algorithm 1** Upper-Confidence-Bound-Network (UCB-Network)

---

Each user in  $G$  follows UCB-user policy

**UCB-user policy for a user  $v$ :**

**Initialization:** For  $1 \leq t \leq K$

- play arm  $t$

**Loop:** For  $K \leq t \leq n$

-  $a^v(t+1) = \underset{j}{\operatorname{argmax}} \hat{\mu}_{m_j^v(t)} + \sqrt{\frac{2 \ln t}{m_j^v(t)}}$

---

where  $\Delta_i = \mu^* - \mu_i$ ,  $\beta \in (0.25, 1)$ ,  $b = m \left[ \frac{2}{4\beta-1} + \frac{2}{(4\beta-1)^2 \ln(1/\beta)} \right] \left( \sum_{j=1}^K \Delta_j \right)$ , and

$C_G$  is a network dependent parameter, defined as follows.

(ii) Let  $\gamma_k = \min\{t \in \{1, \dots, n\} : |\{v \in V : m_i^v(t) \geq l_i = \frac{8 \ln n}{\Delta_i}\}| \geq k\}$  denote the smallest time index when at least  $k$  nodes have access to at least  $l_i$  samples of arm  $i$ . Let  $\eta_k$  be the index of the 'latest' node to acquire  $l_i$  samples of arm  $i$  at  $\gamma_k$ , such that  $\eta_k \neq \eta_{k'}$  for  $1 \leq k, k' \leq m$ . Define  $z_k = T_i(\gamma_k) := (T_i^1(\gamma_k), \dots, T_i^m(\gamma_k))$ , which contains the arm  $i$  counts of all nodes at time  $\gamma_k$ . Then,  $C_{G, l_i}$  is the solution of the following optimisation problem:

$$\begin{aligned} \max \quad & \|z_m\|_1 \\ \text{s.t.} \quad & \exists \text{ a sequence } \{z_k\}_{k=1}^m \\ & z_j(\eta_k) = z_k(\eta_k) \quad \forall j \geq k \\ & \langle z_k, A(\eta_k, \cdot) \rangle \geq l_i, \quad 1 \leq k \leq m \end{aligned} \tag{2}$$

**Remark 1** We propose the result in red, which is a correction of the result from [1]. After checking the proof, we think that we can deduce this tighter upper bound on the expected total regret of  $G$ .

**Corollary 1** For an  $m$ -node star network:

$$\mathbb{E}[R^G(n)] \leq b + (m-1) \sum_{i: \mu_i < \mu^*} \frac{8 \ln n}{\Delta_i} \tag{3}$$

### 3.2 Follow Your Leader (FYL)

They present a second network policy called Follow Your Leader (FYL) for a generic  $m$ -node network. The policy is based on exploiting high-degree hubs in the graph; for this purpose, we need to define notions of dominating sets and dominating set partitions.

**Definition 3** [Dominating set of a graph] A dominating set  $D$  of a graph  $G = (V, E)$  is a subset of  $V$  such that every node in  $V \setminus D$  is adjacent to atleast one of the nodes in  $D$ .

**Definition 4** [Dominating set partition of a graph] Let  $D$  be a dominating set of  $G$ . A dominating set partition based on  $D$  is obtained by partitioning  $V$  into  $|D|$  components such that each component contains a node in  $D$  and a subset of its one hop neighbors.

The FYL policy for an  $m$ -node generic network is outlined in Algorithm 2. Under the FYL policy, all nodes in the dominating set are called *leaders* and all other nodes as *followers*; the follower nodes follow their leaders while choosing an action in a round. The following theorem presents an upper bound on the expected regret of an  $m$ -node star network which employs the FYL policy.

**Theorem 2 (FYL regret bound, star networks)** Suppose the star network  $G$  with a dominating set as the center node, follows the FYL policy to learn a stochastic MAB problem with  $K$  arms. Assume that the rewards lie in  $[0, 1]$ . Then,

$$\mathbb{E}[R^G(n)] \leq d + \sum_{i: \mu_i < \mu^*}^K \frac{8 \ln n}{\Delta_i},$$

where  $d = \left\lceil 2m - 1 + \frac{2m}{4\beta-1} \left( 1 + \frac{1}{(4\beta-1) \ln(1/\beta)} \right) \right\rceil \left( \sum_{j=1}^K \Delta_j \right)$ ,  $\Delta_i = \mu^* - \mu_i$  and  $\beta \in (0.25, 1)$ .

---

**Algorithm 2** Follow Your Leader (FYL) Policy

---

**Input:** Graph  $G$ , a dominating set  $D$  and a dominating set partition

**Leader - Each node in  $D$ :**

Follows the UCB-user policy by using the samples of itself and its neighbors

**Follower - Each node in  $V \setminus D$ :**

In round  $t = 1$ :

- Chooses an action randomly from  $\mathcal{K}$

In round  $t > 1$ :

- Chooses the action taken by the leader in its component, in the previous round  $(t - 1)$
- 

A key insight obtained from Theorem 2 is that an  $m$ -node star network with the FYL policy incurs an expected regret that is lower by a factor  $(m - 1)$ , as compared to UCB-Network.

### 3.3 Lower bounds on the expected regret

We decided not to focus on the lower bounds on the expected regret during this review, but it's important to notice that they prove in [1] that the regret upper bound under the FYL policy meets a lower bound universal to all policies. Hence, they conclude that the FYL policy is order optimal for star networks.

## 4 Our theoretical contributions

## 5 Experimental results

### 5.1 Reproducing experiments from [1]

Considering this is a review project, we decided to first reproduce the experimental results from the paper.

#### 5.1.1 Experiment 1

Performance comparison of UCB-Network policy on various 10 node networks: 2 arms, Bernoulli rewards with means 0.5 and 0.7.




Figure 1: Results for experiment 1 from [1] (100 sample paths).




Figure 2: Our results for experiment 1 (1000 sample paths).

We observe that we obtain exactly the same results as in [1], except for the expected regret of the network at rounds  $> 10000$ . We observe a noisier

#### 5.1.2 Experiment 2

Performance comparison of UCB-Network policy on various 20 node networks: 10 arms, Bernoulli rewards with means  $0.1, 0.2, \dots, 1$ .




Figure 3: Results for experiment 2 from [1] (100 sample paths).




Figure 4: Our results for experiment 2 (1000 sample paths).

### 5.1.3 Experiment 3

Performance comparison of UCB-Network and FYL policies on various star networks: 2 arms, Bernoulli rewards with means 0.5 and 0.7.




Figure 5: Results for experiment 3 from [1] (100 sample paths).




Figure 6: Our results for experiment 3 (1000 sample paths).

## 5.2 Distributed star network

## 6 Conclusion

## References

- [1] Ravi Kumar Kolla, Krishna P. Jagannathan, and Aditya Gopalan. Stochastic bandits on a social network: Collaborative learning with local information sharing. *CoRR*, abs/1602.08886, 2016.

## 7 Citations, figures, tables, references

These instructions apply to everyone.

### 7.1 Citations within the text

The natbib package will be loaded for you by default. Citations may be author/year or numeric, as long as you maintain internal consistency. As to the format of the references themselves, any style is acceptable as long as it is used consistently.

The documentation for natbib may be found at

<http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf>

Of note is the command `\citet`, which produces citations appropriate for use in inline text. For example,

```
\citet{hasselmo} investigated\dotso
```

produces

Hasselmo, et al. (1995) investigated...

If you wish to load the natbib package with options, you may add the following before loading the neurips\_2018 package:

```
\PassOptionsToPackage{options}{natbib}
```

If natbib clashes with another package you load, you can add the optional argument nonatbib when loading the style file:

```
\usepackage[nonatbib]{neurips_2018}
```

As submission is double blind, refer to your own published work in the third person. That is, use “In the previous work of Jones et al. [4],” not “In our previous work [4].” If you cite your other papers that are not widely available (e.g., a journal paper under review), use anonymous author names in the citation, e.g., an author of the form “A. Anonymous.”

### 7.2 Tables

All tables must be centered, neat, clean and legible. The table number and title always appear before the table. See Table 1.

Place one line space before the table title, one line space after the table title, and one line space after the table. The table title must be lower case (except for first word and proper nouns); tables are numbered consecutively.

Note that publication-quality tables *do not contain vertical rules*. We strongly suggest the use of the booktabs package, which allows for typesetting high-quality, professional tables:

<https://www.ctan.org/pkg/booktabs>

This package was used to typeset Table 1.



Table 1: Sample table title

Part		
Name	Description	Size ( $\mu\text{m}$ )
Dendrite	Input terminal	$\sim 100$
Axon	Output terminal	$\sim 10$
Soma	Cell body	up to $10^6$